

# Cosine Similarity

# How to define a good similarity measure?

- Euclidean distance

- $dist(q, d) =$


- $$\sqrt{\sum_{t \in V} [tf(t, q)idf(t) - tf(t, d)idf(t)]^2}$$

- Longer documents will be penalized by the extra words

- We care more about how these two vectors are overlapped

# Why distance is a bad idea

- The Euclidean distance between  $\vec{q}$  and  $\vec{d}_2$  is large even though the distribution of terms in the query  $\vec{q}$  and the distribution of terms in the document  $\vec{d}_2$  are very similar.




# Use angle instead of distance

- Thought experiment: take a document  $d$  and append it to itself. Call this document  $d'$ .
- “Semantically”  $d$  and  $d'$  have the same content
- The Euclidean distance between the two documents can be quite large
- The angle between the two documents is 0, corresponding to maximal similarity.
- Key idea: Rank documents according to angle with query.

# From distance to angle


- Angle: how vectors are overlapped
  - Cosine similarity – projection of one vector onto another



# From angles to cosines

- The following two notions are equivalent.
  - Rank documents in decreasing order of the angle between query and document
  - Rank documents in increasing order of  $\cos(\text{angle}(\text{query}, \text{document}))$
- Cosine is a monotonically decreasing function for the interval  $[0^\circ, 180^\circ]$

# From angles to cosines



- But how – *and why* – should we be computing cosines?

# From angles to cosines

- Cosine of 0 degrees = 1
- Cosine of 90 degree = 0
- Cosine of 180 degree = -1

Cosine of angles varies between -1 to 1.



# Question

Is it a problem for tf-idf that the cosine of the angle between vectors is negative between  $90^\circ$  and  $180^\circ$ ?


- A. Yes, because we don't know that to do with negative scores.
- B. No, because a vector  $d_1$  pointing in the opposite direction of a vector  $d_2$  would just have the same content.
- C. No, because we can always normalize by  $(1 + \cos(d_1, d_2)) / 2$  which is always between 0 and 1.
- D. No, because the cosine of the angles between tf-idf vectors will always be non-negative.

# Answer

- Correct Option is D
- Explanation
  - The cosine of the angles between tf-idf vectors will always be non-negative (i.e. are never more than 90° apart). This is because all the elements of any tf-idf vector are non-negative (i.e. they are in  $\mathbf{R}^n_+$ , where  $n$  is the number of unique terms), and no two vectors in the non-negative orthant are more than 90° apart.

# Cosine similarity

- Angle between two vectors
  - $\text{cosine}(V_q, V_d) = \frac{V_q \times V_d}{|V_q|_2 \times |V_d|_2} = \boxed{\frac{V_q}{|V_q|_2}} \times \frac{V_d}{|V_d|_2}$ 
    - TF-IDF vector
    - Unit vector
  - Document length normalized



# cosine(query,document)

Dot product

Unit vectors

Can we ignore  
query length  
in this  
formula?

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \bullet \frac{\vec{d}}{|\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

$q_i$  is the tf-idf weight of term  $i$  in the query  $d_i$  is the tf-idf weight of term  $i$  in the document

$\cos(\vec{q}, \vec{d})$  is the cosine similarity of  $\vec{q}$  and  $\vec{d}$  ... or, equivalently, the cosine of the angle between  $\vec{q}$  and  $\vec{d}$ .


# Cosine for length-normalized vectors

- For length-normalized vectors, cosine similarity is simply the dot product (or scalar product):

$$\cos(\vec{q}, \vec{d}) = \vec{q} \bullet \vec{d} = \sum_{i=1}^{|V|} q_i d_i$$

for q, d length-normalized.

# Cosine similarity illustrated



# Cosine similarity amongst 3 documents

term	SaS	PaP	WH
affection	115	58	20
jealous	10	7	11
gossip	2	0	6
wuthering	0	0	38

## Term frequencies (counts)

How similar are the novels

**SaS**: *Sense and Sensibility*

**PaP**: *Pride and Prejudice*

**WH**: *Wuthering Heights*?

Note: To simplify this example, we don't do idf weighting.

# 3 documents example contd.

## Log frequency weighting

term	SaS	PaP	WH
affection	3.06	2.76	2.30
jealous	2.00	1.85	2.04
gossip	1.30	0	1.78
wuthering	0	0	2.58

## After length normalization

term	SaS	PaP	WH
affection	0.789	0.832	0.524
jealous	0.515	0.555	0.465
gossip	0.335	0	0.405
wuthering	0	0	0.588

$\cos(\text{SaS}, \text{PaP}) \approx$

$0.789 \times 0.832 + 0.515 \times 0.555 + 0.335 \times 0.0 + 0.0 \times 0.0$

$\approx 0.94$

$\cos(\text{SaS}, \text{WH}) \approx 0.79$

$\cos(\text{PaP}, \text{WH}) \approx 0.69$

Why do we have  $\cos(\text{SaS}, \text{PaP}) > \cos(\text{SaS}, \text{WH})$ ?



# tf-idf example

Document: *car insurance auto insurance*

Query: *best car insurance*

Term	Query		Document
	tf-row	df	tf-row
Auto	0	5000	1
best	1	50000	0
car	1	10000	1
insurance	1	1000	2

Total documents = 1 million

# tf-idf example

Document: *car insurance auto insurance*

Query: *best car insurance*

Term	Query					Document		Product
	tf-raw	tf-wt	df	idf	tf*idf	tf-raw	tf-wt	
auto	0	0	5000	2.3	0	1	1	0
best	1	1	50000	1.3	1.3	0	0	0
car	1	1	10000	2.0	2.0	1	1	2
insurance	1	1	1000	3.0	3.0	2	1.3	3.9

Doc length =  $\sqrt{1^2 + 0^2 + 1^2 + 1.3^2} \approx 1.92$  Note: We have not used query length here so the similarity score is not between 0 and 1

Score =  $(0+0+2+3.9) / (1.92) = 5.9/1.92 = 3.07$

# IDF can be multiplied to only one vector

- $(Tf_{best, q} * \mathbf{IDF}_{best}) * Tf_{best, d} + (Tf_{car, q} * \mathbf{IDF}_{car}) * Tf_{car, d}$
- $Tf_{best, q} * \mathbf{IDF}_{best} * Tf_{best, d} * \mathbf{IDF}_{best} + Tf_{car, q} * \mathbf{IDF}_{car} * Tf_{car, d} * \mathbf{IDF}_{car}$

$$= Tf_{best, q} * Tf_{best, d} * (\mathbf{IDF}_{best})^2 + Tf_{car, q} * Tf_{car, d} * (\mathbf{IDF}_{car})^2$$

IDF of a word is same in query and document so multiplying it with any one of the vectors is enough

# Question

- What is  $N$ , the number of documents, for this example?

(Hint: Note that the document frequency for “car” is 10,000 and its inverse document frequency is 2.0.)

- A. 1 million
- B. 10,000
- C. 500,000
- D. No idea...

# Answer

- Option A is correct
- Explanation:
  - $N=1$  million because we see that, from the formula for idf,  $\log_{10} N / 10,000 = 2.0$  so we have  $N = 10,000 \times 10^2 = 10,000 \times 100 = 10^6$ , which is 1 million.

# Fast computation of cosine in retrieval

- $\text{cosine}(V_q, V_d) = V_q \times \frac{V_d}{|V_d|_2}$ 
  - $|V_q|_2$  would be the same for all candidate docs
  - Normalization of  $V_d$  can be done in index time
  - Only count  $t \in q \cap d$
  - Score accumulator for each query term when intersecting postings from inverted index

# Computing cosine scores

COSINESCORE( $q$ )

```
1  float Scores[ $N$ ] = 0
2  float Length[ $N$ ]
3  for each query term  $t$ 
4  do calculate  $w_{t,q}$  and fetch postings list for  $t$ 
5      for each pair( $d, tf_{t,d}$ ) in postings list
6      do  $Scores[d] + = w_{t,d} \times w_{t,q}$ 
7  Read the array Length
8  for each  $d$ 
9  do  $Scores[d] = Scores[d] / Length[d]$ 
10 return Top  $K$  components of Scores[]
```

# Fast computation of cosine in retrieval

- Maintain a score accumulator for each doc when scanning the postings

Query = “info security”

$S(d,q)=g(t_1)+...+g(t_n)$  [sum of TF of matched terms]

Info: (d1, 3), (d2, 4), (d3, 1), (d4, 5)

Security: (d2, 3), (d4, 1), (d5, 3)

Can be easily applied to TF-IDF weighting!

Accumulators:		d1	d2	d3	d4	d5
info	(d1,3) =>	<b>3</b>	0	0	0	0
	(d2,4) =>	3	4	0	0	0
	(d3,1) =>	3	4	<b>1</b>	0	0
	(d4,5) =>	3	4	1	5	0
security	(d2,3) =>	3	<b>7</b>	1	5	0
	(d4,1) =>	3	7	1	<b>6</b>	0
	(d5,3) =>	3	7	1	6	<b>3</b>



# Advantages of VS Model

- Empirically effective! (Top TREC performance)
- Intuitive
- Easy to implement
- Well-studied/Mostly evaluated
- The Smart system
  - Developed at Cornell: 1960-1999
  - Still widely used
- **Warning: Many variants of TF-IDF!**

# Disadvantages of VS Model

- Assume term independence
- Assume query and document to be the same
- Lack of “predictive adequacy”
  - Arbitrary term weighting
  - Arbitrary similarity measure

# What you should know

- Document ranking v.s. selection
- Basic idea of vector space model
- Two important heuristics in VS model
  - TF
  - IDF
- Similarity measure for VS model